

Smit Thakare

(551) 344-4551 | smit.thakare@columbia.edu | linkedin.com/in/smit-thakare | smitthakare.com

EDUCATION

Columbia University

M.S. in Computer Science, Machine Learning Track

New York, NY
Aug 2025 – Dec 2026

K.J. Somaiya College of Engineering

B.Tech in Information Technology (Honours in AI) — GPA: 9.3 / 10.0

Mumbai, IN
Jul 2021 – May 2025

EXPERIENCE

S&P Global

AI Product Management Intern, Corporate Actions

New York, NY
Jun 2026 – Aug 2026

- Embedded with a business team to translate an ambiguous, manual workflow into a production generative AI solution (LangGraph) on Google Cloud (GKE, Kubeflow) — reading source data from Excel and an existing database to draft structured, reviewable outputs instead of a blank-page manual process.
- Owned the solution end-to-end — data ingestion, entity resolution via fuzzy-matching, and LLM-based classification of each record (MATCH / INSERT / ORPHAN / DISAGREEMENT) — with every output auditable and nothing finalized without a human review step, iterating directly with the business team that used it daily.

Columbia University — Natural Language Processing Lab

Machine Learning Research Assistant — Prof. Kathleen McKeown

New York, NY
Feb 2026 – Present

- Designed and built a six-stage data pipeline to construct a multilingual Hindi dataset from raw source text (~46K sentences, 97.8% Devanagari), enabling structured LLM evaluation on the MASIVE / ARNI benchmarks for a low-resource language.
- Ran continual-learning experiments (mT5-large, Qwen2.5-3B) across a five-language sequence, demonstrating minimal catastrophic forgetting and measurable forward transfer — direct experience productionizing and evaluating LLMs under real-world constraints.

Columbia Business School — Behavioral Research Lab

Data Analyst

New York, NY
Jan 2026 – Present

- Designed and built an automated Python data pipeline (Arctic Shift API) that extracted and structured 200K+ Reddit posts, handling pagination, rate limiting, and deduplication to produce clean, analysis-ready datasets.
- Partnered directly with faculty stakeholders to translate ambiguous research questions into concrete data requirements, delivering reproducible datasets that powered an end-to-end research study.

Savic Technologies

Software Engineering Intern (Full-Stack + ML)

Mumbai, IN
Jan 2025 – Apr 2025

- Owned an ML feature end-to-end — model, API, and frontend — automating resume-parsing with a DistilBERT + T5 pipeline that improved recruitment-screening efficiency by 34%.
- Shipped HR analytics dashboards (Python/Flask REST APIs, SAP integration) for live business users, lifting dashboard engagement 30%.

PROJECTS

LLM Router — Cost-Optimized AI Inference System | Python, DistilBERT, Llama-3.1-8B

- Built a DistilBERT binary classifier that routes each query to a lightweight or large model by predicted difficulty, cutting inference cost 61.6% while retaining 78.7% of large-model accuracy — a production cost/quality tradeoff directly relevant to deploying LLMs at scale.
- Validated the cost–quality tradeoff across 800+ benchmark samples using an evaluation framework built on hypothesis testing, cross-validation, and feature analysis (statistically significant, $p < 0.01$).

Yield Curve Visualizer — Macro Sentiment Analysis | FastAPI, PostgreSQL, Streamlit, Python

- Built a backend data pipeline processing 200+ Treasury/FRED time series into a PostgreSQL warehouse, tuning the schema with caching and composite indexes for 40% faster historical queries.
- Layered an NLP sentiment pipeline over 100+ Fed speeches and used server-side aggregation to render 200+ series in a responsive analytics dashboard for end users.

Questscape — iOS GPS Multiplayer Game | Swift, SwiftUI, Real-Time GPS Multiplayer

- Shipped a live GPS-based multiplayer game to the App Store with a 3-person team (Columbia Startup Studio); grew to 70+ downloads and 43 active users — end-to-end ownership from build to live users.

PUBLICATIONS & CERTIFICATIONS

- Co-author, “MOSAIC-EMO: Multilingual LLMs’ Understanding of Nuanced Affect Across Cultures” — under review, EAACL 2026.
- Extractive text-summarization pipeline (85% accuracy) — Springer ICDMAI 2024.
- Capital IQ Pro — Academy & Tech Support (Basic) Certifications, S&P Global (Jul 2026).

SKILLS

- Generative AI & Agentic Engineering:** Python, LangGraph, LangChain, multi-agent systems, RAG, prompt engineering, Claude/Anthropic & OpenAI APIs, agentic frameworks, production LLM deployment
- Structured Generation & Human-in-the-Loop Systems:** Grounded/retrieval-based generation, output validation, auditable & reviewable pipelines, model evaluation frameworks
- Backend & Data:** FastAPI, Flask, SQL, PostgreSQL, SQLAlchemy, ETL/data pipelines, Google Cloud Platform (GKE, Kubeflow)
- Delivery:** Requirements discovery with business/product stakeholders, cross-functional collaboration, Agile, end-to-end ownership from prototype to production